

Ilia Karmanov

Zurich, Switzerland · ilia.uk@gmail.com · +41 76 803 1861
ilkarman.github.io · github.com/ilkarman · [Google Scholar](#)
UK and EU citizen (no visa sponsorship needed)

Overview

Senior Staff Research Scientist at NVIDIA ADLR, working on multimodal architecture, pre-training, post-training, and evaluation for NVIDIA's Nemotron models. First author of Eclair (shipped open-source as Nemotron Parse), a document-understanding model that generated 50B+ tokens for LLM pre-training and underpins NVIDIA's VLM document understanding. Earlier, published at NeurIPS, ICCV, and BMVC at Qualcomm AI Research. Came to machine learning in 2016 from economics: an MSc at LSE focused on optimal control and causal inference, and development-economics research at the International Growth Centre (Oxford and LSE).

Experience

NVIDIA ADLR, Zurich

Senior Staff Research Scientist

Mar 2025 – present

- Made the case for RL post-training on Nemotron VL and drove a first GRPO proof-of-concept with the systems team, showing it improved the model; this seeded the team's standard RL pipeline
- Improved RL training by inspecting evaluation failures and turning them into better rewards and data; on MMLongBench-Doc this made us the top model under 10B (as of mid-2026)
- Exploring multi-teacher on-policy distillation: letting a smaller student learn answer-correctness from higher-capacity teachers without inheriting a reasoning style it cannot support

Staff Research Scientist

Sep 2022 – Feb 2025

- Contributed to multimodal pre-training and architecture design for Nemotron (LLM and VLM)
- First author of Eclair, a document-understanding model that extracts formatted text, tables, formulas, and reading order. Shipped open-source as Nemotron Parse and adopted across NVIDIA's pipelines for pre-training data (50B+ tokens) and VLM distillation; it underpinned NVIDIA's Llama Nemotron Nano VL, which topped OCRBench v2 on release
- Led the model's design: an asymmetric architecture (larger vision encoder, lightweight decoder), no positional embeddings in the decoder, multi-token prediction to reduce decoding steps, and a single prompt-controlled output format spanning text, layout, and semantic classes

Qualcomm AI Research, Amsterdam

Staff Research Scientist

Dec 2021 – Jul 2022

- Developed hardware-efficient architectures for 4K video-to-video processing (denoising on device)
- Efficient Mixture-of-Experts architecture for conditional computation (BMVC 2022)

Senior Research Scientist

Feb 2020 – Nov 2021

- Created the first weakly-supervised approach for 3D person-localisation on video and WiFi data, published at NeurIPS 2021 (with Max Welling) and IEEE GLOBECOM 2021 (first author)
- Video recognition via motion-augmented self-training (ICCV 2021)

Microsoft, London

Senior Applied Scientist

Dec 2018 – Jan 2020

- Applied computer vision across industry problems: self- and fully-supervised segmentation (medical and seismic imaging), object detection and counting, and real-time video action recognition
- Built synthetic-data pipelines to improve model generalisation

Applied Scientist

Jun 2016 – Nov 2018

- Collaborated with MSR on Z3 theorem-prover for plane-gate allocations (Mixed Integer Linear Programming)
- Created DeepLearningFrameworks: benchmarking deep learning frameworks across single-node and multi-node on Azure. 1,700+ GitHub stars, contributions from framework creators

Charles River Associates, London

Consultant

May 2014 – Jun 2016

- Econometrics for antitrust and M&A cases: regression specification and robustness, mathematical programming, and geospatial algorithms
- Initiated the first data-scientist role within the department

International Growth Centre, University of Oxford

Research Economist

Oct 2013 – May 2014

- Applied econometrics for development economics at the IGC, the Oxford and LSE research centre funded by the UK government, under the firm-capabilities theme (large firms, entrepreneurship, trade) directed by Prof. Paul Collier
- Performed fieldwork in India, Pakistan, and Uganda with local researchers, presenting findings to policymakers

Publications

I Karmanov, A Deshmukh, L Voegtler, P Fischer, K Chumachenko, T Roman, J Seppanen, J Parmar, J Jennings, A Tao, K Sapra. Eclair: Extracting Content and Layout with Integrated Reading Order for Documents. arXiv 2025

K Chumachenko, A Deshmukh, J Seppanen, **I Karmanov**, C Chen, L Voegtler, P Fischer, et al. Nemotron Parse 1.1. arXiv 2025

Contributing author on NVIDIA tech reports: Nemotron 3 Nano Omni (arXiv 2026), Nemotron Nano V2 VL (arXiv 2025), and Nemotron-H Hybrid Mamba-Transformer (arXiv 2025)

Z Li, G Chen, S Liu, ... **I Karmanov**, L Voegtler, P Fischer, ... Z Yu. Eagle 2: Building Post-Training Data Strategies from Scratch for Frontier Vision-Language Models. NeurIPS 2025

A Royer, **I Karmanov**, A Skliar, B Ehteshami Bejnordi, T Blankevoort. Revisiting Single-gated Mixtures of Experts. BMVC 2022

FG Zanjani, **I Karmanov**, H Ackermann, D Dijkman, S Merlin, M Welling, F Porikli. Modality-Agnostic Topology Aware Localization. NeurIPS 2021

I Karmanov, FG Zanjani, S Merlin, I Kadampot, D Dijkman. WiCluster: Passive Indoor 2D/3D Positioning using WiFi without Precise Labels. IEEE GLOBECOM 2021

K Gavriluk, M Jain, **I Karmanov**, CGM Snoek. Motion-Augmented Self-Training for Video Recognition at Smaller Scale. ICCV 2021

Y Ren, J Lu, A Beletchi, Y Huang, **I Karmanov**, et al. Hand Gesture Recognition using 802.11ad mmWave. IEEE WCNC 2021

Book: M Salvaris, D Dean, WH Tok, **I Karmanov**. Deep Learning with Azure. Apress, 2018

J Costa-Font, F Cowell (with research contributions from **I Karmanov**). European Identity and Redistributive Preferences. CESifo Working Paper, 2015

Patents

12 US patent applications filed at Qualcomm (several granted), across RF positioning, gesture recognition, and efficient architectures.

Service

Reviewer for NeurIPS, ICLR, and ICML.

Education

London School of Economics

MSc Economics

2012 – 2013

- Thesis: Optimal Control Analysis for Tax Avoidance. Modelling a firm's optimisation problem in an infinite horizon setting, introducing reputation as an accumulated asset, solved with Hamiltonian in an optimal-control setting. Advisor: Prof. Frank Cowell

BSc Economics, First Class Honours

2009 – 2012